

Cell-MPI

Mastering the Cell Broadband Engine architecture through a Boost based parallel communication library



Sebastian Schaetz, Joel Falcou,
Lionel Lacassagne


Digiteo Foundation, LRI - University
Paris South XI, CEA LIST

May 17, 2011


Why a talk about a library for the Cell in 2011

- Heterogeneous and composable architectures are not uncommon, powerful and worth studying.
- We present useful concepts that apply to all of them.
- We illustrate the lessons we learned as we used Boost libraries on a constricted platform and
- elaborate what choices we had to make and why we made them as we created a Boost-like library for this platform.

Cell Broadband Engine - Schematic



A similar architecture - Multi-GPU Schematic



Cell Broadband Engine - The good stuff

- Power architecture core paired with up to 8 streamlined vector co-processors: 204.8 GFlops/s (single) 102.4 GFlops/s (double)
- High data transfer bandwidth: theoretical 204.8 GB/s
- Good performance/watt (0.87 double precision GFlops/s per Watt for IBM BladeCenter QS22)

Due to these advantages, the CBE is a good fit for multimedia and vector processing applications as well as scientific computation.

Cell Broadband Engine - The bad stuff

- Distributed system on one chip, explicit communication necessary
- SPE Memory limitations
 - 256kB for code and data per SPE
 - no overflow detection
- Communication intricacies
 - packet size
 - address alignment
 - explicit DMA
- Optimization for speed
 - SIMD (assembler-like)
 - convoluted pipeline mechanism

Due to these restrictions, the complexity of programming the CBE is comparable to writing code for embedded systems.

Writing code for the CBE

- PPE and SPE entry points in separate main functions
- Compilers: ppu-gcc, spu-gcc
- SPE object file passed to ppu-embedspu to generate library exports symbol that is accessible from PPE code
- PPE creates thread for each SPE and loads the symbol
- Argument passed to thread is accessible in SPE through argument vector
- Usual approach: argument is pointer to structure in main memory; structure is loaded to SPE through explicit DMA call:

```
1 | /* DMA control block information from system memory. */  
2 | mfc_get((void*)&parms, parm_ptr, (sizeof(parms)+15)&~0xF, tag, td, rd);  
3 | mfc_write_tag_mask(1<<tag);  
4 | mfc_read_tag_status_all(); /* Wait for DMA to complete */
```

Writing code for the CBE - continued

- "Getting started" can be tedious when developing for the CELL since compilation procedure and startup are not trivial

- CMake to the rescue:
 - Great tool to simplify basically any build-related steps
 - Find all required libraries and binaries on the system
 - Low-level macros: `ACTIVATE_PPE_COMPILER()`, `ACTIVATE_SPE_COMPILER()`
 - `ADD_SPE_MODULE(target symbol file0 file1 ... fileN)`

- C++ and Boost to the rescue:
 - Wrap recurring boilerplate code in clearly laid out functions and classes
 - A kernel function should be declared and behave like a free function

Cell-MPI Bootstrapping

- Launching a kernel function passing a data structure

```
struct mydatastruct { int x; int y; int z; };
```

- Kernel is defined with:

```
1 | BEGIN_CELL_KERNEL()  
2 | {  
3 |     mydatastruct * ptr;  
4 |     SPE_Custom(ptr);  
5 |     RETURN((ptr->x + ptr->y) * ptr->z);  
6 | }  
7 | END_CELL_KERNEL()
```

- In PPE code the kernel is registered with
PPE_REGISTER_KERNEL(kernel);
- The runtime is initialized with PPE_Init();

Cell-MPI Bootstrapping - continued

- The kernel is then called asynchronously:


```
1 | mydatastruct mydata(1, 5, 7);  
2 | PPE_Run(kernel, mydata);
```

- The PPE can wait for kernel completion: `PPE_Sync()`;
- and access the kernels return value:

```
1 | int returnvalues[CBE_MPI_NUM_SPE];  
2 | PPE_Return(&returnvalues[0]);
```

- The runtime is finalized with `PPE_Finalize()`;

Cell-MPI Bootstrapping Mechanism



Cell-MPI Bootstrapping - Boostified

- A kernel can be declared in both PPE and SPE code with:

```
1 | SPE_FUNCTION(kernel_, kernel, (int x) (int y) (int z) );
```

- and implemented as a free function in SPE code

```
1 | int kernel(int x, int y, int z)  
2 | {  
3 |     return (x+y) * z;  
4 | }
```

- It can then be called as a free function from PPE code:

```
1 | int * returnvalues = kernel(2, 5, 7);
```

- or asynchronously:

```
1 | kernel_async(2, 5, 7);  
2 | PPE_Sync();
```

So we do C++ but...

The architecture forces some restrictions especially on the SPE part of the library:

- Compilation without run-time type information
- No dynamic memory allocation for predictable footprint
- Custom, lightweight STL compatible allocators
- Exception handling deactivated

Exception emulation

Due to architecture limitations we emulate exceptions:

- An exception stops the kernel and notifies the PPE
- Only an error code is "thrown":

```
1 | #define THROW(errno) {spe_errno = errno ; SPE_Finalize(-1); exit(0);}
```

- The PPE translates the error code into real exceptions:

```
1 | struct spe_error_bundle  
2 | { std::vector<spe_error_data> exception_info; };
```

```
1 | typedef boost::error_info<struct tag_spe_error_info_bundle,  
2 |     spe_error_bundle> spe_error_info_bundle;
```

```
1 | struct spe_runtime_exception : virtual boost::exception {};
```

Exception emulation - continued

- If desired SPE exceptions can interrupt PPE execution
- To define errors we use the same trick as in boosted bootstrapping:

```

1 | ERROR(MPI_TAG_MISMATCH, 7, "Send receive tag mismatch")
2 | ERROR(BOOST_FUNCTION_BAD_CALL, 12, "Bad boost function call")
3 | ERROR(BAD_ALLOC, 14, "bad alloc")

```

- Compiled with the SPE compiler (`#ifdef _SPE_`) generates:

```

1 | enum { MPI_TAG_MISMATCH = 7, BOOST_FUNCTION_BAD_CALL = 12,
2 |        BAD_ALLOC = 14 };

```

- And with the PPE compiler generates a vector of objects:

```

1 | struct spe_error_struct
2 | { int id; const char * symbol; const char * message; };

```

Unit Testing

- Boost.Test is great but builds don't fit SPEs:
libboost_unit_test_framework.so.1.45.0: 998kB
- First idea: boost/detail/lightweight_test.hpp
misses a lot of the Boost.Test goodness

Enter SPE-Unit:

- Compromise between lightweight and feature-complete
- Designed after Boost.Test

Unit Testing - SPE Unit Features

- Only one test suite is available: `CBE_MPI_SPEUNIT_AUTO_TEST_SUITE()`;
- Tests are started explicitly:

```
1 | uint32_t result = CBE_MPI_SPEUNIT_RUN_TEST_SUITE();  
2 | SET_RETURN_VALUE(result);
```

- The powerful `AUTO_TEST_CASE_TEMPLATE(testname, T, typelist)` and a normal template `TEST_CASE_TEMPLATE(testname)` are included
- Different test tool levels are supported: `WARN_*`, `CHECK_*`, `REQUIRE_*`
- Strings can be disabled to reduce overhead (silent mode)
- Emulated SPE exceptions can be validated with test tools like `CBE_MPI_REQUIRE_THROW`

Unit Testing - Example

```
1 typedef boost::mpl::vector_c<int,1,2,4,8,16> aligned_alloc_alignments;
2
3 CBE_MPI_SPEUNIT_AUTO_TEST_SUITE();
4
5 CBE_MPI_SPEUNIT_AUTO_TEST_CASE_TEMPLATE( aligned_malloc_free_test, T,
6     aligned_alloc_alignments )
7 {
8     aligned_ptr<void,T::value> ptr = aligned_malloc<T::value>(T::value);
9     CBE_MPI_SPEUNIT_REQUIRE_EQUAL(is_aligned<T::value>(ptr.get()),true);
10    cbe_mpi::aligned_free(ptr);
11    CBE_MPI_SPEUNIT_REQUIRE_EQUAL(ptr.get(),((void*)(0)));
12 }
13
14 int kernel(void)
15 {
16     uint32_t result = CBE_MPI_SPEUNIT_RUN_TEST_SUITE();
17     SET_RETURN_VALUE(result);
18 }
```

Data Transfer - Single Buffer

```

ii = in.get();
oo = out.get();

for(int i=0; i<iterations; i++) {



    spe_ppe_get_c(in.get(), cd->inbuf1+(SPE_Rank()+i*SPE_Size()*slicesize*sizeof(float),
    slicesize_padded*sizeof(float)); load

    harris_simd(ii, oo, cd->slice_dimx, cd->slice_dimy, 0, PADY, buf1.get(), buf2.get(), buf3.get()); calc

    spe_ppe_put_c(cd->outbuf1+(SPE_Rank()+i*SPE_Size()*slicesize*sizeof(float) +
    (cd->slice_dimx*PADY)*sizeof(float), oo, slicesize*sizeof(float)); store
}

```

Data Transfer



Data Transfer - Double Buffering

```
spe_ppe_get_async_c(in1.get(), cd->inbuf1+SPE_Rank()*slicesize*sizeof(float), slicesize_padded*sizeof(float), 9);
```

```
for(int i=0; i<iterations; i++) {
  if(i%2 == 0) {
    spe_ppe_get_async_c(in2.get(), cd->inbuf1+(SPE_Rank()+i+1)*SPE_Size()*slicesize*sizeof(float),
      slicesize_padded*sizeof(float), 10);

    dma_synchronize_c(9); dma_synchronize_c(11);
    ii = in1.get(); oo = out1.get();
  } else {
    spe_ppe_get_async_c(in1.get(), cd->inbuf1+(SPE_Rank()+i+1)*SPE_Size()*slicesize*sizeof(float),
      slicesize_padded*sizeof(float), 9);

    dma_synchronize_c(10); dma_synchronize_c(12);
    ii = in2.get(); oo = out2.get();
  }
}
```

lead-in

load,
sync

```
harris_simd(ii, oo, cd->slice_dimx, cd->slice_dimy, 0, PADY, buf1.get(), buf2.get(), buf3.get());
```

calc

```
if(i%2 == 0) {
  spe_ppe_put_async_c(cd->outbuf1+(SPE_Rank()+i*SPE_Size())* slicesize*sizeof(float) +
    (cd->slice_dimx*PADY)*sizeof(float), out1.get(), slicesize*sizeof(float), 11);
} else {
  spe_ppe_put_async_c(cd->outbuf1+(SPE_Rank()+i*SPE_Size())* slicesize*sizeof(float) +
    (cd->slice_dimx*PADY)*sizeof(float), out2.get(), slicesize*sizeof(float), 12);
}
}
```

store

```
spe_ppe_put_c(cd->outbuf1 + (SPE_Rank()+((iterations-1)*SPE_Size())) * slicesize*sizeof(float) +
  (cd->slice_dimx*PADY)*sizeof(float), oo, slicesize*sizeof(float));
```

lead-out

Double Buffering - Operations - Input Segment

- Start loading first segment (lead-in)
`operator =()`
- Start loading next segment
`operator ++(int)`
- Wait for segment to be ready for computation
`operator *()`
- Signal that computation on current segment is finished
`operator ++(int)`
- Check if end of data is reached
`operator ==()`


Double Buffered Segmented Input Iterator

```
1  template<typename T> struct remote_segmented_input_iterator
2  {
3      // allocate required buffers
4      remote_segmented_input_iterator(...) {}
5
6      // start loading first buffer
7      void operator= (const addr64 & base_address_) { }
8
9      // wait for current segment to arrive and return pointer to it
10     T* operator *() {}
11
12     // start loading new data and increment current segment
13     inline void operator++(int) {}
14
15     // check if iterator has reached a position
16     bool operator ==(const addr64 & b) const {}
17 };
```


Double Buffered Segmented Iterator Example

```
1 remote_segmented_input_iterator<float> it(depth,  
2   ssize, slicer(ssize));  
3 remote_segmented_output_iterator<float> ot(depth,  
4   ssize, slicer(ssize));  
5  
6 for(it = input, ot = output; /* lead-in */  
7   it!=input+overall_size; /* check end */  
8   it++, ot++) // load next, store current  
9 {  
10  float * in = *it; float * out = *ot; // synchronize  
11  harris_simd(in, out, cd->slice_dimx, cd->slice_dimy,  
12  0, PADY, buf1.get(), buf2.get(), buf3.get());  
13 }
```

Double Buffered Segmented Iterator - Slicer



Double Buffered Segmented Iterator - Slicer



Multi-Buffered Segmented Iterator - Features

- `remote_vector<T>` for more expressive code:


```
1 // PPE:
2 std::vector<float> v(1024*1024); kernel(v);
3 // SPE:
4 kernel(remote_vector<float> v) {
5     remote_segmented_input_iterator<T> it(depth, ssize, slicer(ssize));
6     for(it = v.begin(); it!=v.end(); it++) {
7         float * in = *it;
8         /* computation */
9     }
10 }
```

- Read, write- and read-write Iterators with minimum buffer depth of 3
- Various slicers

2D Multi-Buffered Segmented Iterator

- Native 2D data transfer support through DMA lists
- Difference to regular iterator:
 - Slice size is 2D
 - Supports `remote_vector_2D`
 - Slicer takes 2D arguments:


```
slicer_2D(size_2d_t vector_dim, size_2d_t slice_dim);
```
- Ideal for image processing:




High-Level Inter-SPE Communication: MPI

- Interprocess communication by message passing, SPEs send and receive message
- API specification, used in high performance computing


Features:

- Virtual topology of processes
- Synchronization
- Point to point communication
- Collective communication


MPI Collectives - Broadcast




MPI Collectives - Scatter and Gather



MPI Collectives - Reduce and All to All



$$X = Op(A_1, Op(Op(B_1, Op(C_1, D_1))))$$



MPI Interface - Example

```
1 communicator world;
2
3 if (world.rank() == 0)
4 {
5     char s1[] = "Hello";
6     world.send(1, 0, s1, sizeof(s1));
7     char s2[6];
8     world.recv(1, 1, s2, sizeof(s2));
9 }
10 else if (world.rank() == 1)
11 {
12     char s1[6];
13     world.recv(0, 0, s1, sizeof(s1));
14     char s2[] = "world";
15     world.send(0, 1, s2, sizeof(s2));
16 }
17 // Hello world from SPE 0, Hello world from SPE 1
```

MPI Interface - Communicator

```
1 class communicator
2 {
3     void barrier();
4
5     template <typename T> void send(int dst, int tag, const T& value);
6     template <typename T> void send(int dst, int tag, const T* values, int n);
7     template <typename T> request isend(int dst, int tag, const T& value);
8     ...
9     template <typename T> status recv(int source, int tag, T& value);
10    template <typename T> status recv(int source, int tag, T* values, int n);
11    template <typename T> request irectv(int source, int tag, T& value);
12    ...
13    communicator include(uint16_t first, uint16_t last);
14    communicator exclude(uint16_t first, uint16_t last);
15    friend bool operator==(const communicator& c1, const communicator& c2);
16 };
```


MPI Interface - Request and Status

```
1 // represents current request
2 class request
3 {
4     request() {};
5     status wait();
6     boost::optional<status> test();
7 };
8
9 // represents status of a request
10 class status
11 {
12     int32_t source() const;
13     int32_t tag() const;
14     int32_t error() const;
15 };
```


MPI Interface - Collectives Interface

```
1  template<typename T, typename Op>
2  void reduce(const communicator & comm, const T & in,
3             T & out, Op op, int root);
4
5  template<typename T, typename Op>
6  void reduce(const communicator & comm, const T & in,
7             Op op, int root);
8
9  template<typename T, typename Op>
10 void reduce(const communicator & comm, const T * in,
11            int n, T * out, Op op, int root);
12
13 template<typename T, typename Op>
14 void reduce(const communicator & comm, const T * in,
15            int n, Op op, int root);
```

MPI Header



MPI Protocol



MPI Types

We don't do Boost.Serialization but

- you may register your POD type:

```
1 | struct gps_position { /* POD */ };  
2 | namespace cbe_mpi  
3 | {  
4 |     CBE_MPI_USER_POD_DATATYPE(gps_position);  
5 | }
```

- or you may specialize send/receive methods:

```
1 | template <typename T>  
2 | request isend(cbe_mpi::communicator & comm, int dst,  
3 |     int tag, T data, int n);  
4 |  
5 | template <typename T>  
6 | request irecv(cbe_mpi::communicator & comm, int src,  
7 |     int tag, T data, int n);
```

Registering POD Types


How we identify your type:

```
1  template<typename T>
2  struct cbe_mpi_user_pod_type_id { static void get() {} };
3
4  #define CBE_MPI_USER_POD_DATATYPE(CppType) \
5  template<> \
6  struct is_mpi_datatype< CppType > \
7  : boost::mpl::bool_<true> {}; \
8      \
9  inline int get_mpi_datatype(const CppType &) \
10 { \
11     return 0x80000000 | \
12         (int)&cbe_mpi_user_pod_type_id< CppType >::get(); \
13 }
```

Sending std::vector

```
1  template <typename T>
2  request isend(cbe_mpi::communicator com,
3              int dest, int tag, const std::vector<T> * values, int)
4  {
5      int vectorsize = values->size();
6      com.send(dest, tag, &vectorsize, 1);
7      return com.isend(dest, tag, &(*values)[0], vectorsize);
8  }
9
10 template <typename T>
11 request irecv(cbe_mpi::communicator com,
12             int source, int tag, std::vector<T> * values, int)
13 {
14     int vectorsize;
15     com.recv(source, tag, &vectorsize, 1);
16     values->resize(vectorsize);
17     return com.irecv(source, tag, &(*values)[0], vectorsize);
18 }
```

MPI - Sending Unaligned Data



Conclusion

- Build process can be simplified with CMake
- Boilerplate code can be simplified with the help of Boost (e.g. PP)
- Ambiguity of functions or macros in different compilation units can be exploited

- Optimal Boost solutions have to be adapted to fit embedded architecture
- Sweet spot between generic code and efficiency must be found

- Difficult low-level code can be wrapped nicely in C++ Concepts
- C++ Concepts can be even more powerful on special purpose hardware

Thank you for you kind attention.